

# Automatic Differentiation

CSE 849 Deep Learning  
Spring 2025

Zijun Cui

- Let's first finish the backpropagation

- Scalar functions of Affine functions

$$\begin{array}{ccc}
 & & d_z \times d_y \\
 D = f(\mathbf{z}) & & \mathbf{z} = \mathbf{W}\mathbf{y} + \mathbf{b} \\
 \text{scalar} & & \begin{array}{ccc} d_z \times 1 & d_y \times 1 & d_z \times 1 \end{array}
 \end{array}$$

Matching Dimension:

$$\begin{array}{ccc}
 & & d_z \times d_y \\
 \nabla_{\mathbf{y}} D = \nabla_{\mathbf{z}}(D) \mathbf{W} & & \\
 \begin{array}{ccc} 1 \times d_y & 1 \times d_z & \end{array}
 \end{array}$$

$$\begin{array}{ccc}
 \nabla_{\mathbf{b}} D = \nabla_{\mathbf{z}}(D) & & \\
 \begin{array}{ccc} 1 \times d_z & 1 \times d_z & \end{array}
 \end{array}$$

$$\begin{array}{ccc}
 & & d_y \times d_z \\
 \nabla_{\mathbf{W}} D = \mathbf{y} \nabla_{\mathbf{z}}(D) & & \\
 \begin{array}{ccc} d_y \times 1 & 1 \times d_z & \end{array}
 \end{array}$$

$$\begin{array}{ccc}
 \text{Why? } \Delta D = \alpha \Delta \mathbf{y} & & \\
 \text{scalar} & d_y \times 1 & 
 \end{array}$$

$$\begin{array}{ccc}
 \Delta D = \alpha \Delta \mathbf{W} & & \\
 \text{scalar} & d_z \times d_y & 
 \end{array}$$

So  $\alpha$  must be  $1 \times d_y$

and  $\alpha = \nabla_{\mathbf{y}} D$

and  $\alpha = \nabla_{\mathbf{W}} D$  with dimension  $d_y \times d_z$

**Wrong.....**

- Scalar functions of Affine functions

$$\Delta D = \sum_{i,j} \frac{\partial D}{\partial W_{ij}} \Delta W_{ij} \quad \text{Frobenius Inner Product for matrix}$$

$$\mathbf{z} = \mathbf{W}\mathbf{y} + \mathbf{b}$$

We want to convert so that  $\mathbf{W}$  becomes variables  
That is why we want to transpose

$$\mathbf{z}^\top = \mathbf{y}^\top \mathbf{W}^\top + \mathbf{b}^\top$$

$$\nabla_{\mathbf{W}^\top} D = \nabla_{\mathbf{z}^\top} D \nabla_{\mathbf{W}^\top} \mathbf{z}^\top = \nabla_{\mathbf{z}^\top} D \mathbf{y}^\top$$


$$\nabla_{\mathbf{W}} D = (\nabla_{\mathbf{W}^\top} D)^\top = \mathbf{y} \nabla_{\mathbf{z}} D$$

# Special Case: Application to a network

- Scalar functions of Affine functions

$$\mathbf{z}_k = \mathbf{W}_k \mathbf{y}_{k-1} + \mathbf{b}_k$$


$$Div = Div(\mathbf{z}_k) \quad \Rightarrow \quad \nabla_{\mathbf{y}_{k-1}} Div = \nabla_{\mathbf{z}_k} Div \mathbf{W}_k$$



$$\nabla_{\mathbf{b}_k} Div = \nabla_{\mathbf{z}_k} Div$$

$$\nabla_{\mathbf{W}_k} D = \mathbf{y}_{k-1} \nabla_{\mathbf{z}_k} Div$$

# The backward pass



The divergence is a nested function:  $Div(Y(z_N))$


$$\nabla_{z_N} Div = \underbrace{\nabla_Y Div}_{\text{Already computed}} \cdot \underbrace{\nabla_{z_N} Y}_{\text{New term}} = \nabla_Y Div \cdot J_Y(z_N)$$

Already computed

New term

First compute the derivative of the divergence w.r.t  $Y$ . The actual derivative depends on the divergence function.

# The backward pass



$$\nabla_{y_{N-1}} Div = \nabla_{z_N} Div \cdot \nabla_{y_{N-1}} z_N$$


$$z_N = W_N y_{N-1} + b_N \Rightarrow \nabla_{y_{N-1}} z_N = W_N$$

$$\nabla_{y_{N-1}} Div = \nabla_{z_N} Div \mathbf{W}_N$$

Already computed


New term

# The backward pass





# The backward pass



$$\nabla_{\mathbf{z}_{N-1}} Div = \nabla_{\mathbf{y}_{N-1}} Div J_{\mathbf{y}_{N-1}}(\mathbf{z}_{N-1})$$

The Jacobian will be a diagonal matrix for scalar activations


$$\nabla_{\mathbf{z}_{N-1}} Div$$

$$\nabla_{\mathbf{z}_{N-1}} Div = \nabla_{\mathbf{y}_{N-1}} Div \cdot \nabla_{\mathbf{z}_{N-1}} \mathbf{y}_{N-1}$$

Already computed

New term

# The backward pass




$$\nabla_{y_{N-2}} Div = \nabla_{z_{N-1}} Div \cdot \nabla_{y_{N-2}} z_{N-1}$$

$$\Rightarrow \nabla_{y_{N-2}} Div = \nabla_{z_{N-1}} Div \mathbf{W}_{N-1}$$

$$\nabla_{W_{N-1}} Div = y_{N-2} \nabla_{z_{N-1}} Div$$

$$\nabla_{b_{N-1}} Div = \nabla_{z_{N-1}} Div$$

# The backward pass





$$\nabla_{z_1} Div = \nabla_{y_1} Div J_{y_1}(z_1)$$

$$\nabla_{w_1} Div = x \nabla_{z_1} Div$$

$$\nabla_{b_1} Div = \nabla_{z_1} Div$$

# Setting up for digit recognition



Training data



- Simple Problem: Recognizing “2” or “not 2”
- Single output with sigmoid activation
  - $Y \in (0,1)$
  - $d$  is either 0 or 1
- Use KL divergence
- Backpropagation to compute derivatives
  - To apply in gradient descent to learn network parameters

# Recognizing the digit

Training data

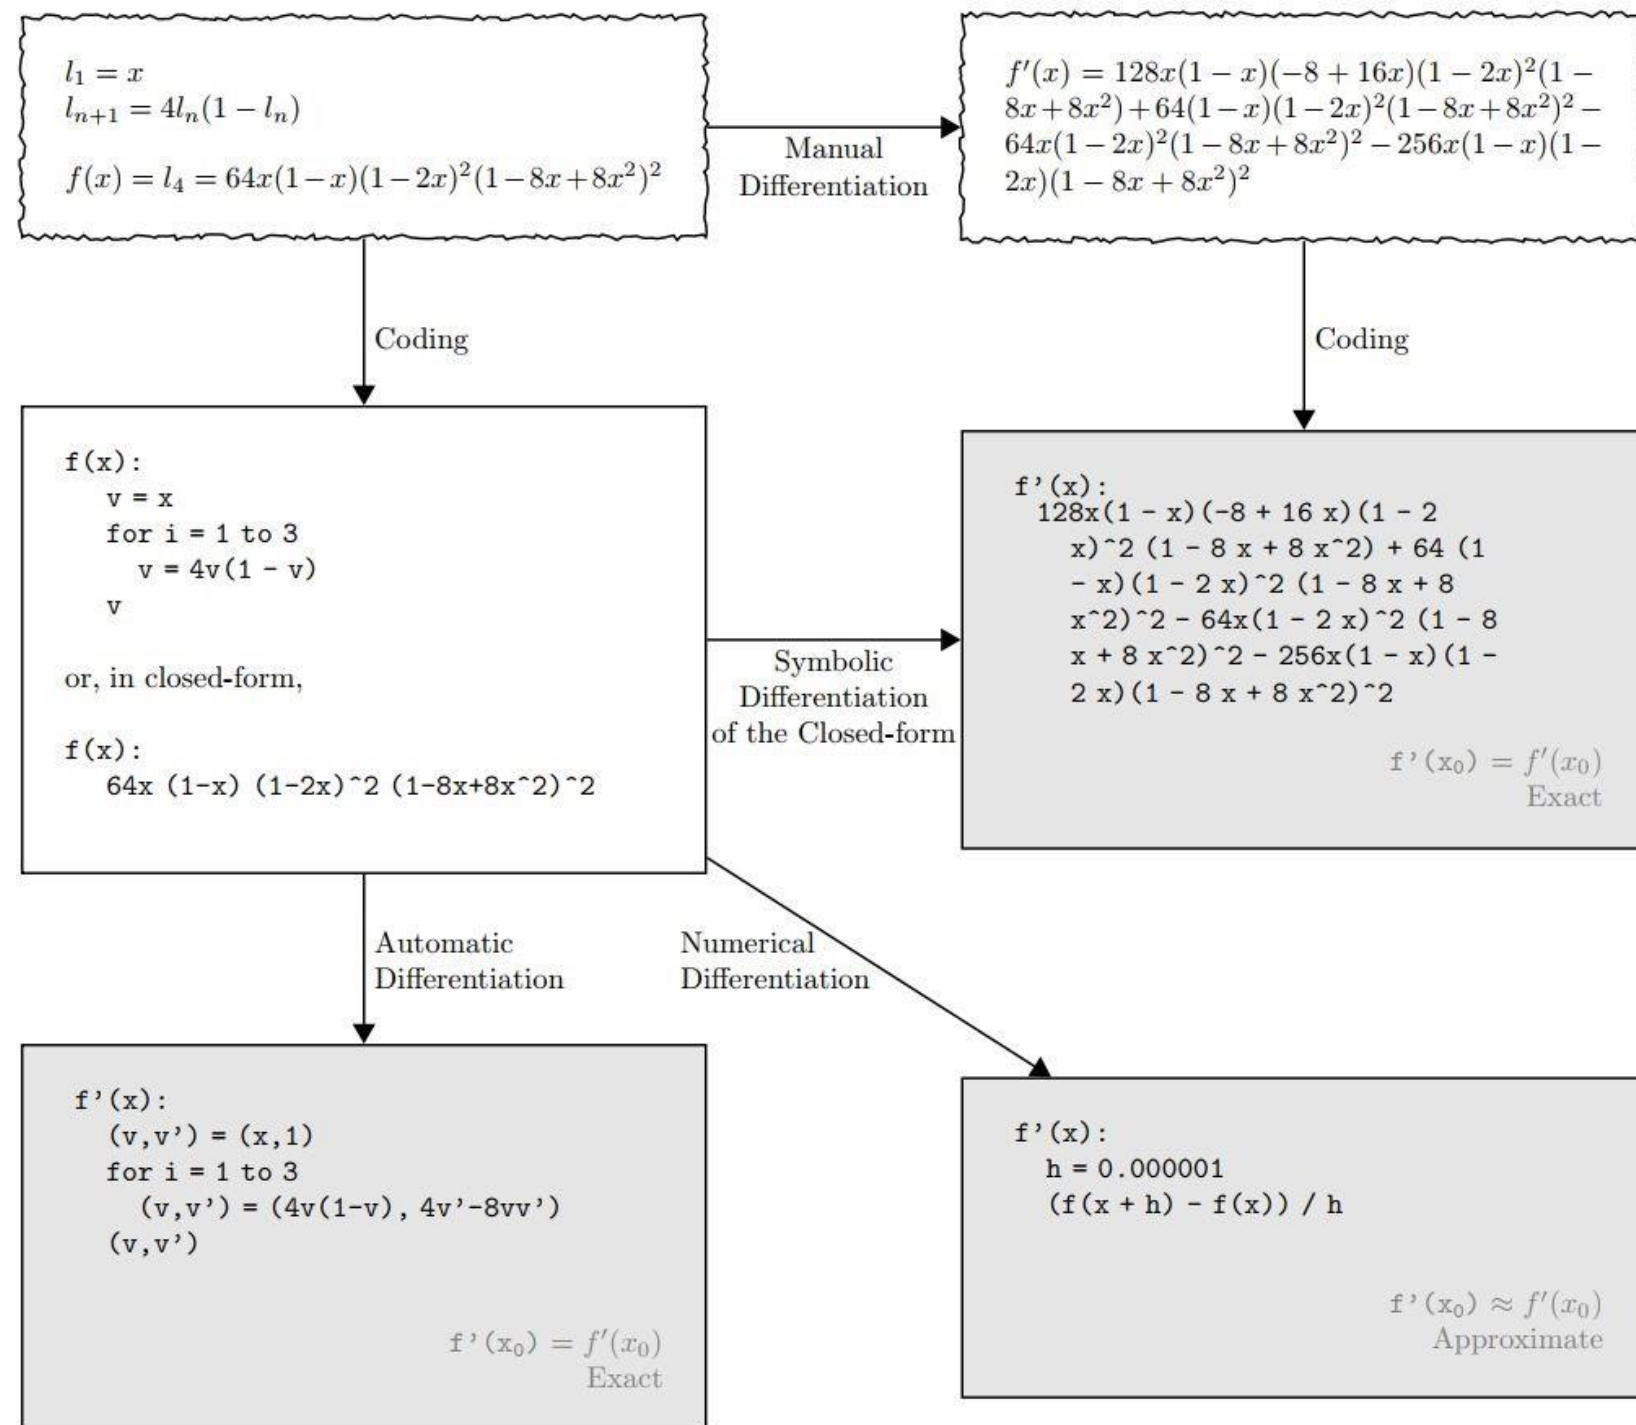


- More complex problem: Recognizing digit
- Network with 10 (or 11) outputs
  - First ten outputs correspond to the ten digits
    - Optional 11th is for none of the above
- Softmax output layer:
  - Ideal output: One of the outputs goes to 1, the others go to 0
- Backpropagation with KL divergence
  - To compute derivatives for gradient descent updates to learn network

- Back to today's topic on Automatic Differentiation

# Derivatives as code

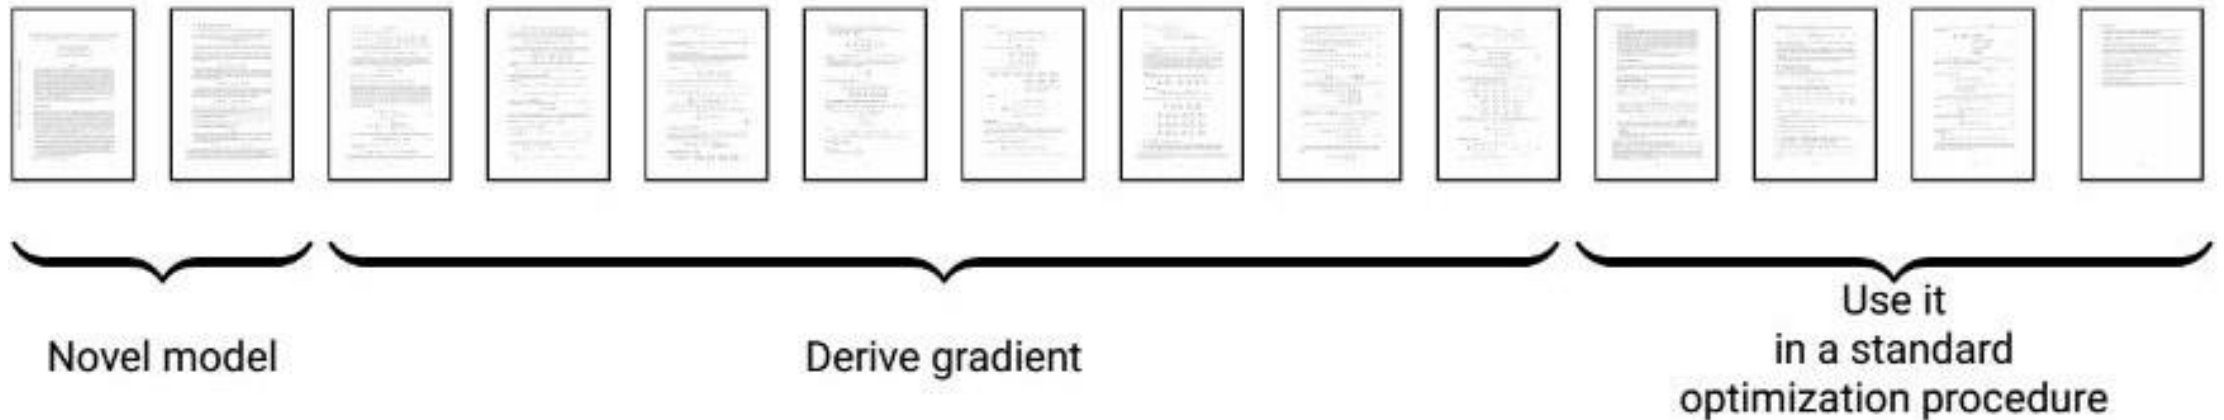
We can compute the derivatives **not just of mathematical functions, but of general programs** (with control flow)



# Manual Differentiation

You can see papers like this:

anisotropic CVT over a sound mathematical framework. In this article a new objective function is defined, and both this function and its gradient are derived in closed-form for surfaces and volumes. This method opens a wide range of possibilities, also described in the



Analytic derivatives are needed for **theoretical insight**

- analytic solutions, proofs
- mathematical analysis, e.g., stability of fixed points

**Unnecessary when we just need derivative evaluations** for optimization



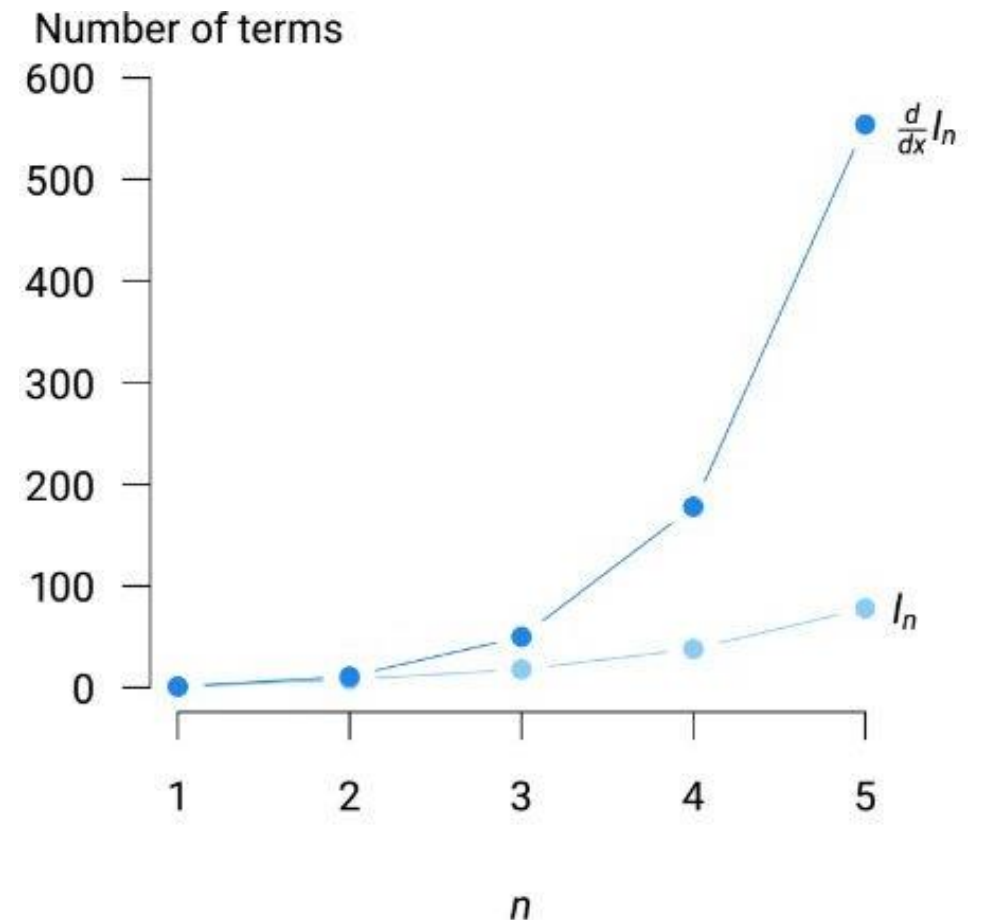
# Symbolic differentiation

Symbolic computation with Mathematica, Maple, Maxima, and deep learning frameworks such as Theano

## Problem: expression swell

Logistic map  $l_{n+1} = 4l_n(1 - l_n)$ ,  $l_1 = x$

$n$	$l_n$	$\frac{d}{dx}l_n$
1	$x$	1
2	$4x(1 - x)$	$4(1 - x) - 4x$
3	$16x(1 - x)(1 - 2x)^2$	$16(1 - x)(1 - 2x)^2 - 16x(1 - 2x)^2 - 64x(1 - x)(1 - 2x)$
4	$64x(1 - x)(1 - 2x)^2(1 - 8x + 8x^2)^2$	$128x(1 - x)(-8 + 16x)(1 - 2x)^2(1 - 8x + 8x^2) + 64(1 - x)(1 - 2x)^2(1 - 8x + 8x^2)^2 - 64x(1 - 2x)^2(1 - 8x + 8x^2)^2 - 256x(1 - x)(1 - 2x)(1 - 8x + 8x^2)^2$



# Symbolic differentiation

- Mathematica's derivatives for one layer of soft ReLU (univariate case):

**D[Log[1 + Exp[w \* x + b]], w]**

Out[11]= 
$$\frac{e^{b+wx} w}{1 + e^{b+wx}}$$

Roger Grosse

- Derivatives for two layers of soft ReLU:

In[19]:= **D[Log[1 + Exp[w2 \* Log[1 + Exp[w1 \* x + b1]] + b2]], w1]**

Out[19]= 
$$\frac{e^{b_1+b_2+w_1 x+w_2 \text{Log}[1+e^{b_1+w_1 x}]} w_2 x}{\left(1 + e^{b_1+w_1 x}\right) \left(1 + e^{b_2+w_2 \text{Log}[1+e^{b_1+w_1 x}]}\right)}$$

# Symbolic differentiation

**Problem:** only applicable to **closed-form mathematical functions**

You can find the derivative of

```
In [1]: def f(x):  
        return 64 * (1-x) * (1-2*x)^2 * (1-8*x+8*x*x)^2
```

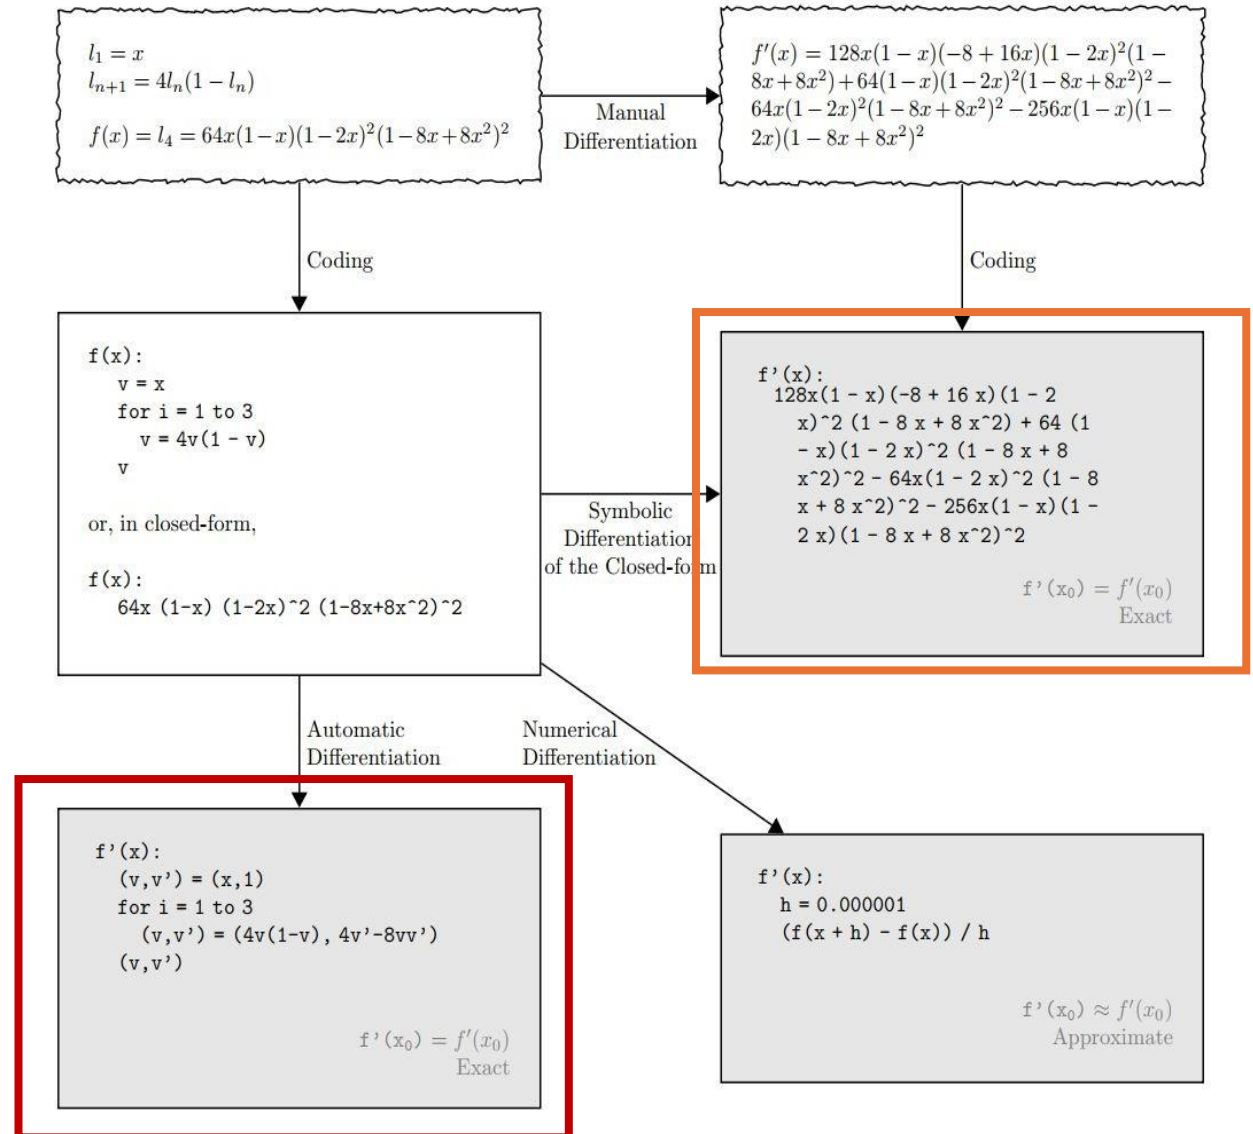
but not of

```
In [2]: def f(x,n):  
        if n == 1:  
            return x  
        else:  
            v = x  
            for i in range(1,n):  
                v = 4*v*(1-v)  
            return v
```

There might not be a convenient formula for the derivatives.

# Autodiff Versus Symbolic differentiation

- The goal of autodiff is not a formula, but a procedure for computing derivatives.



# Numerical differentiation

Finite difference approximation of  $\nabla f$ ,  $f : \mathbb{R}^n \rightarrow \mathbb{R}$

$$\frac{\partial f(\mathbf{x})}{\partial x_i} \approx \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x})}{h}, \quad 0 < h \ll 1$$

**Problem:** needs to be evaluated  $n$  times, once with each  $\mathbf{e}_i \in \mathbb{R}^n$

**Problem:** we must select  $h$  and we face **approximation errors**

# Numerical differentiation

Finite difference approximation of  $\nabla f$ ,  $f : \mathbb{R}^n \rightarrow \mathbb{R}$

$$\frac{\partial f(\mathbf{x})}{\partial x_i} \approx \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x})}{h}, \quad 0 < h \ll 1$$


Better approximations exist:

- Higher-order finite differences

e.g., center difference:

$$\frac{\partial f(\mathbf{x})}{\partial x_i} = \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x} - h\mathbf{e}_i)}{2h} + O(h^2)$$

These increase rapidly in complexity  
and **never completely eliminate the error**



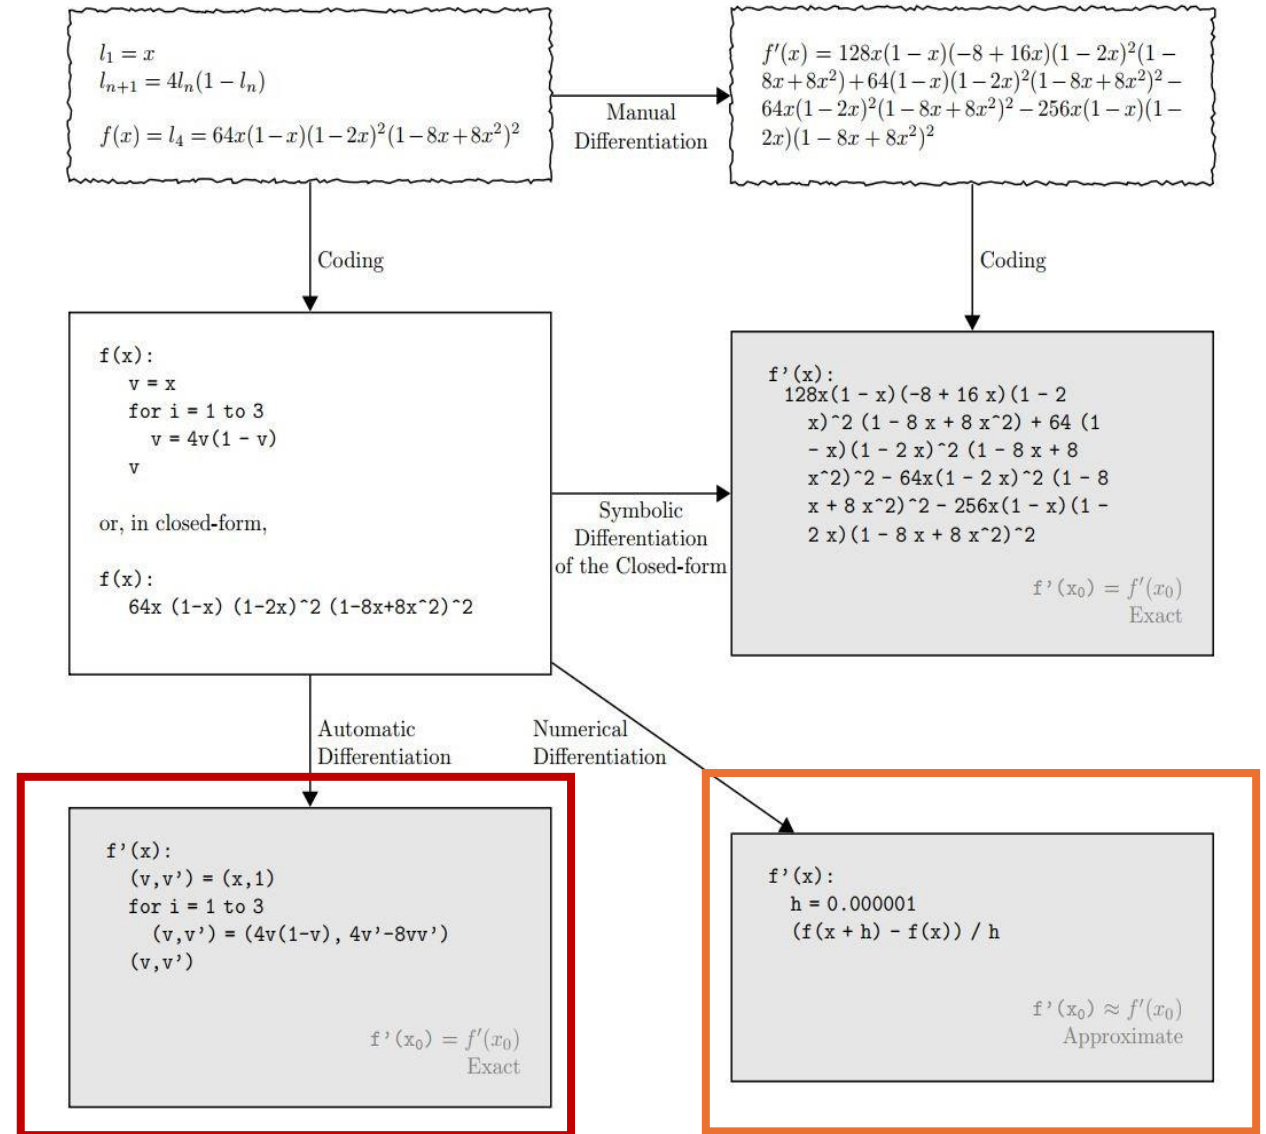
# Autodiff Versus Finite Differences

Finite differences.

Still extremely useful as a **quick check of our gradient implementations**

Normally, we only use it for testing.

Autodiff is both efficient and numerically stable. **Is exact!**



# Automatic differentiation

If we don't need analytic derivative expressions, we can **evaluate a gradient exactly** with only one forward and one reverse execution

$$f : \mathbb{R}^n \rightarrow \mathbb{R} \quad \nabla f(\mathbf{x}) = \left( \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n} \right)$$

In machine learning, this is known as **backpropagation** or “backprop”

- Automatic differentiation is more than backprop
- Or, backprop is a specialized *reverse mode* automatic differentiation



*Nature* 323, 533–536 (9 October 1986)

## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California,  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,  
Pittsburgh, Philadelphia 15213, USA

---

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a



# Confusing Terminology

- **Automatic differentiation (autodiff)** refers to a general way of taking a program which computes a value, and automatically constructing a procedure for computing derivatives of that value.
- **Backpropagation** is the special case of autodiff applied to neural nets  
But in machine learning, we often use backprop synonymously with autodiff
- **Autograd** is the name of a particular autodiff package.  
But lots of people started using “autograd” to mean “autodiff”

# What Autodiff Is

An autodiff system will convert the program into a sequence of **primitive operations** which have specified routines for computing derivatives.

In this representation, backprop can be done in a completely mechanical way.

**Original program:**

$$z = wx + b$$

$$y = \frac{1}{1 + \exp(-z)}$$

$$L = \frac{1}{2}(y - t)^2$$

**Sequence of primitive operations:**

$$t_1 = wx$$

$$z = t_1 + b$$

$$t_3 = -z$$

$$t_4 = \exp(t_3)$$

$$t_5 = 1 + t_4$$

$$y = 1/t_5$$

$$t_6 = y - t$$

$$t_7 = t_6^2$$

$$L = t_7/2$$


# Automatic differentiation

All numerical algorithms, when executed, evaluate to compositions of a finite set of elementary operations with known derivatives

- Called a **trace** or a **Wengert list** (Wengert, 1964)
- Alternatively represented as a **computational graph** showing dependencies

$$f(a, b) = \log(ab)$$

$$\nabla f(a, b) = (1/a, 1/b)$$



$f(a, b)$ :

$c = a * b$

$d = \log(c)$

return  $d$

# Automatic differentiation

Primal: The value computed during the forward pass of a computational graph

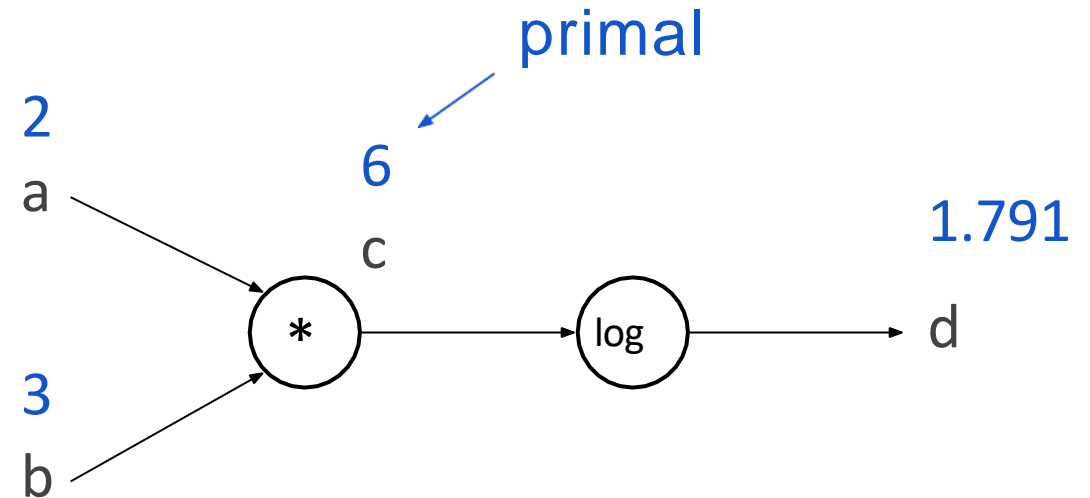
$f(a, b)$ :

$c = a * b$

$d = \log(c)$

return  $d$

$1.791 = f(2, 3)$



# Automatic differentiation

$f(a, b)$ :

$c = a * b$


$d = \log(c)$

return  $d$

$1.791 = f(2, 3)$

$[0.5, 0.333] = f'(2, 3)$


$$\nabla f(a, b) = (1/a, 1/b)$$



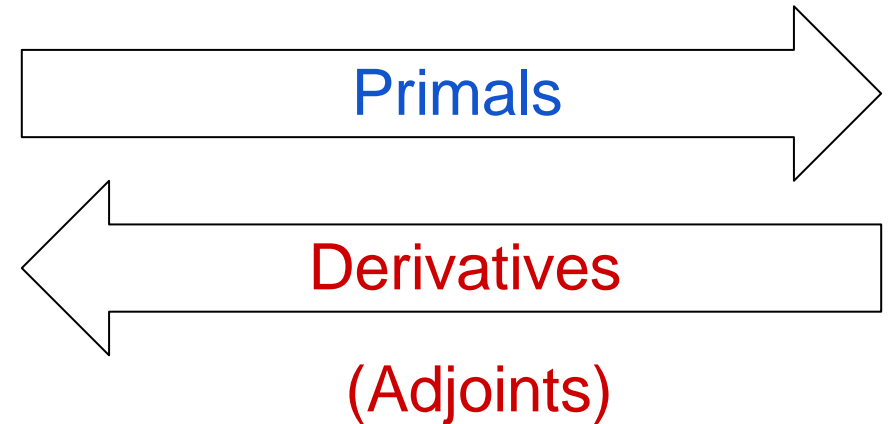
# Automatic differentiation

Two main flavors

**Forward** mode



**Reverse** mode (a.k.a. backprop)



**Nested combinations**

(higher-order derivatives, Hessian–vector products, etc.)

- Forward-on-reverse
- Reverse-on-forward
- ...

# Forward mode

Primals

Derivatives (tangents)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

$f(x1, x2)$ :


$v1 = x1 * x2$

$v2 = \log(x2)$

$y1 = \sin(v1)$

$y2 = v1 + v2$

return  $(y1, y2)$



# Forward mode

Primals

Derivatives (tangents)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

$f(x1, x2)$ :

$v1 = x1 * x2$


$v2 = \log(x2)$

$y1 = \sin(v1)$

$y2 = v1 + v2$

return  $(y1, y2)$

$f(2, 3)$





# Forward mode

Primals

Derivatives (tangents)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

$f(x_1, x_2)$ :

$$v_1 = x_1 * x_2$$


$$v_2 = \log(x_2)$$

$$y_1 = \sin(v_1)$$

$$y_2 = v_1 + v_2$$

return  $(y_1, y_2)$

$f(2, 3)$



$$\frac{\partial x_1}{\partial x_1} = 1$$

# Forward mode

Primals

Derivatives (tangents)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

$f(x_1, x_2)$ :

$v_1 = x_1 * x_2$


$v_2 = \log(x_2)$

$y_1 = \sin(v_1)$

$y_2 = v_1 + v_2$

return  $(y_1, y_2)$

$f(2, 3)$



$$\frac{\partial x_2}{\partial x_1} = 0$$

# Forward mode

Primals

Derivatives (tangents)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

$f(x_1, x_2)$ :

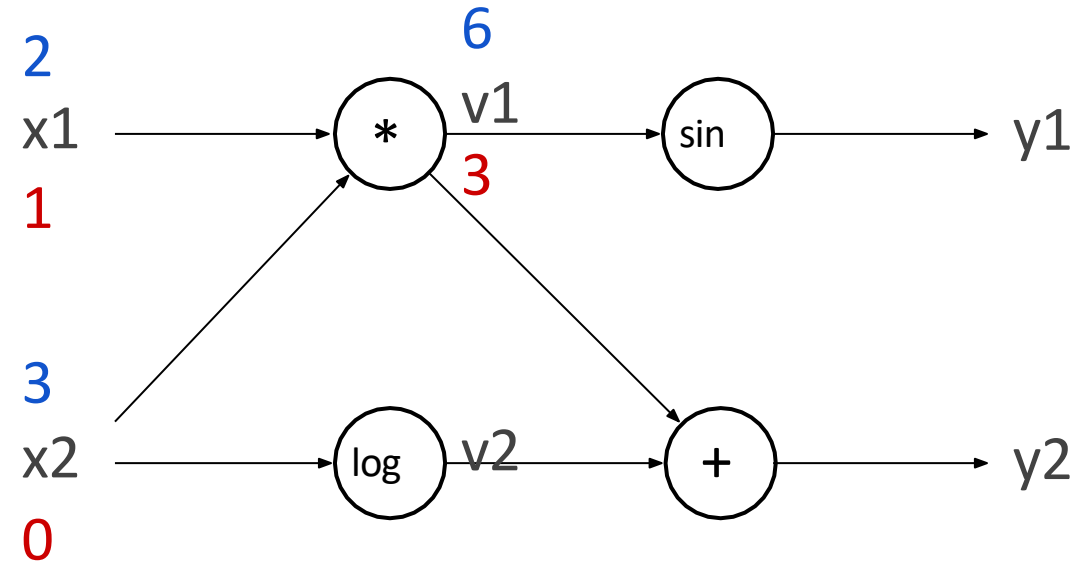
$v_1 = x_1 * x_2$

$v_2 = \log(x_2)$

$y_1 = \sin(v_1)$

$y_2 = v_1 + v_2$

return  $(y_1, y_2)$



$f(2, 3)$

$$\frac{\partial v_1}{\partial x_1} = \frac{\partial x_1}{\partial x_1} x_2 + x_1 \frac{\partial x_2}{\partial x_1} = x_2$$

# Forward mode

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Primals

Derivatives (tangents)

$f(x_1, x_2)$ :

$$v_1 = x_1 * x_2$$


$$v_2 = \log(x_2)$$

$$y_1 = \sin(v_1)$$

$$y_2 = v_1 + v_2$$

return  $(y_1, y_2)$

$f(2, 3)$



$$\frac{\partial v_2}{\partial x_1} = \frac{1}{x_2} \frac{\partial x_2}{\partial x_1} = 0$$

# Forward mode


$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Primals

Derivatives (tangents)

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial x_1} = \cos(v_1) \frac{\partial v_1}{\partial x_1}$$

# Forward mode


$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Primals

Derivatives (tangents)

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_2}{\partial x_1} = \frac{\partial v_1}{\partial x_1} + \frac{\partial v_2}{\partial x_1}$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)





# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial y_1} = 1$$

# Reverse mode


$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Primals

Derivatives (adjoints)

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial y_2} = 0$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial v_1} =$$

# Reverse mode

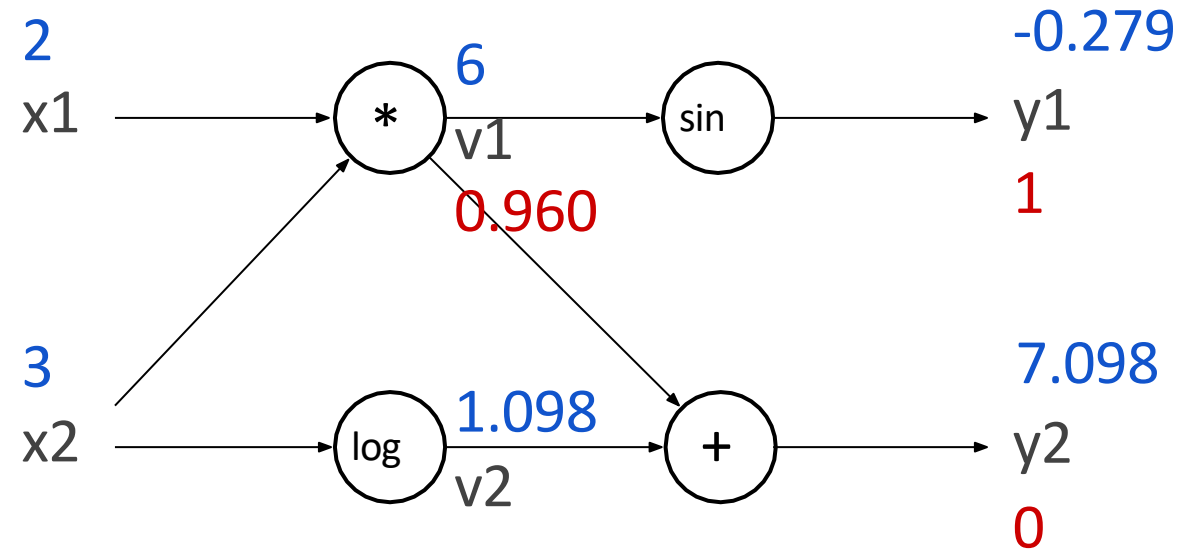
Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial v_1} = \cos(v_1) \frac{\partial y_1}{\partial y_1}$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial v_2} = 0$$



# Reverse mode


$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Primals

Derivatives (adjoints)

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial x_1} =$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial x_1} = \frac{\partial v_1}{\partial x_1} \frac{\partial y_1}{\partial v_1} = x_2 \frac{\partial y_1}{\partial v_1}$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial x_2} =$$

# Reverse mode


Primals

Derivatives (adjoints)

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

```
f(x1, x2):  
  v1 = x1 * x2  
  v2 = log(x2)  
  y1 = sin(v1)  
  y2 = v1 + v2  
  return (y1, y2)
```

f(2, 3)



$$\frac{\partial y_1}{\partial x_2} = \frac{\partial v_1}{\partial x_2} \frac{\partial y_1}{\partial v_1} + \frac{\partial v_2}{\partial x_2} \frac{\partial y_1}{\partial v_2} = x_1 \frac{\partial y_1}{\partial v_1}$$

# Forward vs reverse summary

In the extreme  $\mathbf{f} : \mathbb{R} \rightarrow \mathbb{R}^m$   
use forward mode to evaluate

$$\left(\frac{\partial f_1}{\partial x}, \dots, \frac{\partial f_m}{\partial x}\right)$$

In the extreme  $f : \mathbb{R}^n \rightarrow \mathbb{R}$   
use reverse mode to evaluate

$$\nabla f(\mathbf{x}) = \left(\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n}\right)$$

In general  $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$  the Jacobian  $\mathbf{J}_f(\mathbf{x}) \in \mathbb{R}^{m \times n}$  can be evaluated in

- $O(n \text{ time}(\mathbf{f}))$  with forward mode
- $O(m \text{ time}(\mathbf{f}))$  with reverse mode

Reverse performs better when  $n \gg m$

# Autograd

```
import autograd.numpy as np ← very sneaky!
from autograd import grad

def sigmoid(x):
    return 0.5*(np.tanh(x) + 1)

def logistic_predictions(weights, inputs):
    # Outputs probability of a label being true according to logistic model.
    return sigmoid(np.dot(inputs, weights))

def training_loss(weights):
    # Training loss is the negative log-likelihood of the training labels.
    preds = logistic_predictions(weights, inputs)
    label_probabilities = preds * targets + (1 - preds) * (1 - targets)
    return -np.sum(np.log(label_probabilities))
```

... (load the data) ...


```
# Define a function that returns gradients of training loss using Autograd.
training_gradient_fun = grad(training_loss)

# Optimize weights using gradient descent.
weights = np.array([0.0, 0.0, 0.0])
print "Initial loss:", training_loss(weights)
for i in xrange(100):
    weights -= training_gradient_fun(weights) * 0.01

print "Trained loss:", training_loss(weights)
```

- The rest of this lecture covers how Autograd is implemented.
- Source code for the original Autograd package:  
<https://github.com/HIPS/autograd>
- Autodidact, a pedagogical implementation of Autograd — you are encouraged to read the code.  
<https://github.com/mattjj/autodidact>

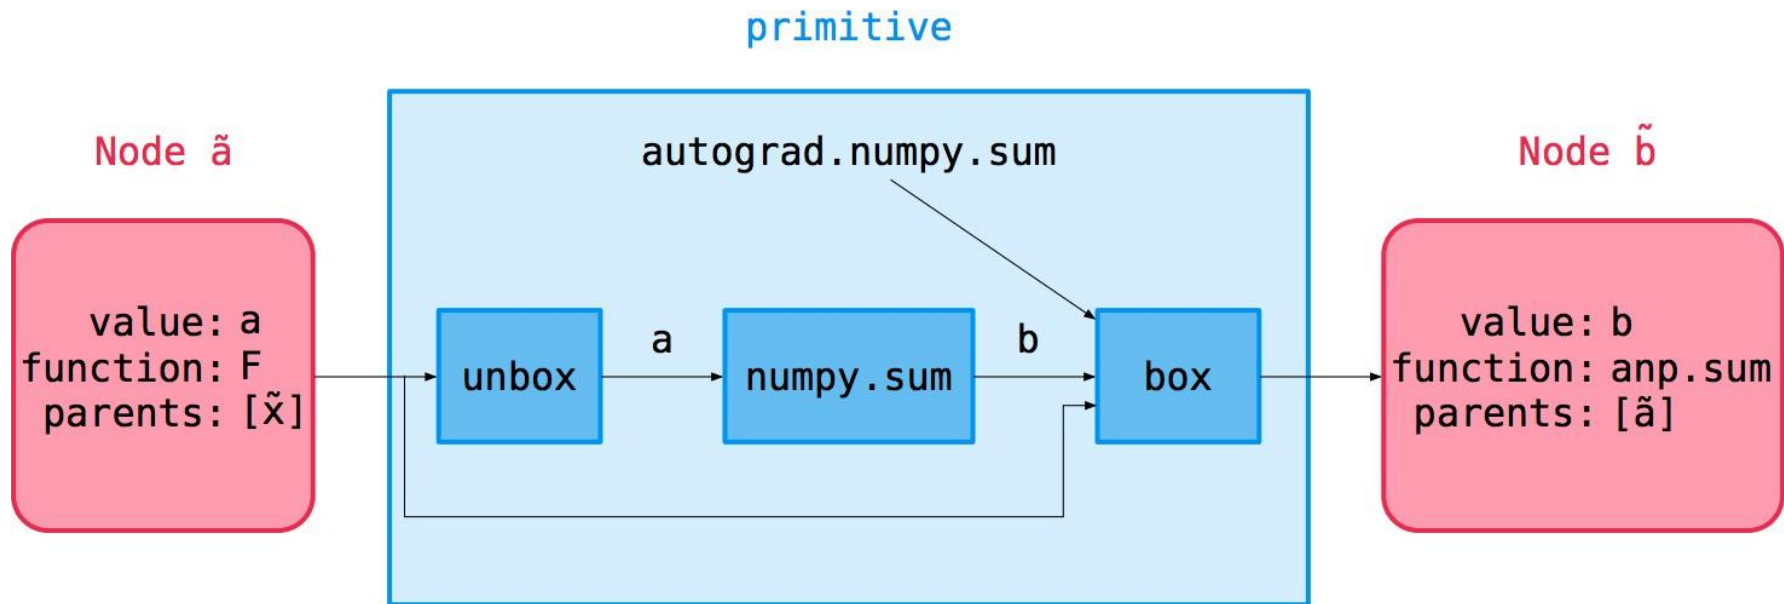
# Building the Computation Graph



- Most autodiff systems, including Autograd, explicitly construct the computation graph.
  - Some frameworks like TensorFlow provide mini-languages for building computation graphs directly. Disadvantage: need to learn a totally new API.
  - Autograd instead builds them by **tracing** the forward pass computation, allowing for an interface nearly indistinguishable from NumPy.
- The **Node** class (defined in `tracer.py`) represents a node of the computation graph. It has attributes:
  - `value`, the actual value computed on a particular set of inputs
  - `fun`, the primitive operation defining the node
  - `args` and `kwargs`, the arguments the op was called with
  - `parents`, the parent Nodes

# Building the Computation Graph

- Autograd's fake NumPy module provides primitive ops which look and feel like NumPy functions, but secretly build the computation graph.
- They wrap around NumPy functions:



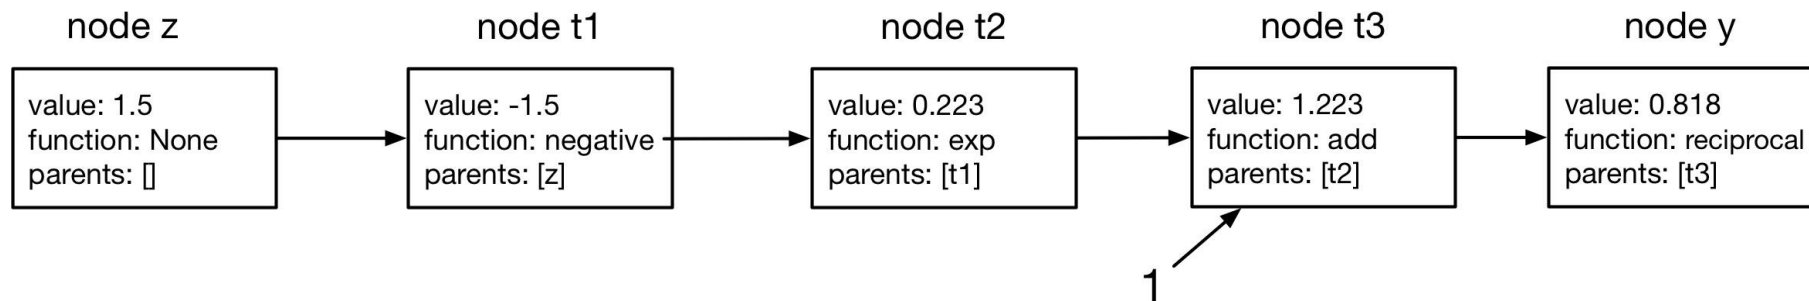


# Building the Computation Graph

Example:

```
def logistic(z):  
    return 1. / (1. + np.exp(-z))  
  
# that is equivalent to:  
def logistic2(z):  
    return np.reciprocal(np.add(1, np.exp(np.negative(z))))  
  
z = 1.5  
y = logistic(z)
```

73 /  
23



# Vector-Jacobian Products

- Implement the primitive operations in vectorized form.

- The Jacobian is the matrix of partial derivatives:

$$\mathbf{J} = \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{matrix} & \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \begin{matrix} \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{matrix} \end{matrix}$$

- The backprop equation (single child node) can be written as a **vector-Jacobian product (VJP)**:

$$\bar{x}_j = \sum_i \bar{y}_i \frac{\partial y_i}{\partial x_j} \qquad \bar{\mathbf{x}} = \bar{\mathbf{y}}^T \mathbf{J}$$

- That gives a row vector 1 by n. We can treat it as a column vector by taking

$$\bar{\mathbf{x}} = \mathbf{J}^T \bar{\mathbf{y}}$$

Note: usually don't explicitly construct the Jacobian. It's simpler and more efficient to compute the VJP directly.

# Vector-Jacobian Products


- For each primitive operation, we must specify VJPs for *each* of its arguments. Consider  $y = \exp(x)$ .
- This is a function which takes in the output gradient (i.e.  $\bar{y}$ ), the answer ( $y$ ), and the arguments ( $x$ ), and returns the input gradient ( $\bar{x}$ )
- `defvjp` (defined in `core.py`) is a convenience routine for registering VJPs. It just adds them to a dict.
- Examples from `numpy/numpy vjps.py`

```
defvjp(negative, lambda g, ans, x: -g)
defvjp(exp,      lambda g, ans, x: ans * g)
defvjp(log,      lambda g, ans, x: g / x)

defvjp(add,      lambda g, ans, x, y : g,
               lambda g, ans, x, y : g)
defvjp(multiply, lambda g, ans, x, y : y * g,
               lambda g, ans, x, y : x * g)
defvjp(subtract, lambda g, ans, x, y : g,
               lambda g, ans, x, y : -g)
```

# Backward Pass

- Backprop computations are more modular if we view them as message passing.



The backwards pass is defined in core.py.

```
def backward_pass(g, end_node):
    outgrads = {end_node: g}
    for node in toposort(end_node):
        outgrad = outgrads.pop(node)
        fun, value, args, kwargs, argnums = node.recipe
        for argnum, parent in zip(argnums, node.parents):
            vjp = primitive_vjps[fun][argnum]
            parent_grad = vjp(outgrad, value, *args, **kwargs)
            outgrads[parent] = add_outgrads(outgrads.get(parent), parent_grad)
    return outgrad

def add_outgrads(prev_g, g):
    if prev_g is None:
        return g
    return prev_g + g
```

# Summary

- We saw three main parts to the code:
  - tracing the forward pass to build the computation graph
  - vector-Jacobian products for primitive ops
  - the backwards pass
- Building the computation graph requires fancy NumPy gymnastics, but other two items are basically what I showed you.
- You're encouraged to read the full code (< 200 lines!) at:  
<https://github.com/mattjj/autodidact/tree/master/autograd>